

Origin of the CEEX–PHOKHARA discrepancy in the pion angular distributions

KLOE large-angle scenario, $e^+e^- \rightarrow \pi^+\pi^-\gamma$ at NLO, $F_\pi = 1$, no vacuum polarisation

Jeremy Paltrinieri*

23 September 2026

Abstract

In the KLOE large-angle (KLOE-LA) scenario the NLO predictions of CEEX and PHOKHARA 10.0 for $e^+e^- \rightarrow \pi^+\pi^-\gamma$ agree in the total cross section and in every charge-even distribution, but disagree in the pion polar angles: the forward–backward asymmetry of the π^+ is $A(\theta^+) = 0.2375$ in CEEX (and 0.2372 in BabaYaga) against 0.2787 in PHOKHARA. We show that the difference is entirely due to two errors in the one-photon ISR×FSR interference at NLO in PHOKHARA (routine `ampL0`, pion branch of the full-NLO mode): a misplaced parenthesis that adds three helicity terms of the tree interference a second time, and a block that keeps the tree interference together with an analytic ISR soft function that the full-NLO mode replaces by the eikonal elsewhere. The second error makes the PHOKHARA asymmetry depend on the unphysical soft-photon cutoff. A point-by-point comparison of the one-photon matrix elements shows that the two errors account for the whole charge-odd difference to 2.5×10^{-5} of the tree. With both corrected, PHOKHARA gives $A(\theta^+) = 0.23724(10)$ averaged over the cutoff and 0.23725(3) with 5×10^9 events, independent of the cutoff; CEEX (0.23725(3)) and BabaYaga (0.23724) agree with it to 3×10^{-5} .

Contents

1	Summary of findings	2
2	Setup	2
3	The discrepancy	3
4	Strategy	4
5	Tests	5
6	The two errors in PHOKHARA	7
7	Result with the fix	8
8	Caveats and open points	10
9	Other findings	11
A	Reproduction and files	11

*Analysis carried out with Claude Code (Anthropic). Repository: https://github.com/jpaltrinieri/CEEX_PHOKHARA_ANALYSIS

1 Summary of findings

1. **The disagreement is reproduced** with a fresh PHOKHARA build and fresh CEEX runs (Section 3): totals agree to 10^{-4} , charge-even shapes agree, and only $\theta^\pm$ and θ_{avg} differ ($\chi^2/\text{ndf} \approx 10^3$). At LO the two codes agree exactly.
2. **The PHOKHARA asymmetry depends on the soft cutoff w** : $A(\theta^+)$ changes by $+0.0156$ per decade of w while σ is constant (Section 5.1).
3. **The whole effect sits in the charge-odd part of the one-photon contribution** (Born + virtual + soft); the hard two-photon emission agrees between the codes (Section 5.3).
4. **Two errors in PHOKHARA explain it point by point** (Sections 5.4 and 6):
 - **H1** – `phokhara_10.0.f:9969-9982`: the factor $(v_f + v_s)$ multiplies only the last of four helicity terms of the tree interference;
 - **H2** – block B5, the ISR one-loop current times the FSR tree (`phokhara_10.0.f:9998-10007`; `ISRcurrent_no_mass_terms.f:134,210`): an extra $\text{Re}[F^*I](1 + S_{\text{ISR}}(w))$.

Subtracting both leaves a per-point residual of 2.5×10^{-5} of the tree.

5. **With both corrected**, $A(\theta^+)$ no longer depends on w and agrees with CEEX and BabaYaga, in total and separately in the one- and two-photon parts (Section 7). Totals and charge-even distributions are unchanged.

2 Setup

Scenario. KLOE-LA (Strong2020 KLOE-I), $\sqrt{s} = 1.02$ GeV: $50^\circ \leq \theta_\pm \leq 130^\circ$; $|p_\pm^z| > 90$ MeV or $p_\pm^\perp > 160$ MeV; at least one photon with $50^\circ \leq \theta_\gamma \leq 130^\circ$ and $E_\gamma > 20$ MeV; $0.1 \leq M_{\pi\pi}^2 \leq 0.85$ GeV². Angles are measured from the electron beam in both codes. Histograms: the 18 KLOE-I observables with 600 uniform bins. The forward–backward asymmetry is $A = [\sigma(\theta < 90^\circ) - \sigma(\theta > 90^\circ)]/\sigma$.

Codes.

- **CEEX-main** (commit e05d007): `pipig -scenario=KLOE-LA -NLO`, one-loop amplitudes from GoSam (`eepipig_noVP_1L_onshell`, `eepipigg`), YFS soft subtraction with cutoff $E_{\text{min}} = 10^{-5}$ GeV (10^{-4} GeV only in the cutoff checks); 5 MPI ranks per task.
- **PHOKHARA 10.0** (`PHOKARA_10.0 @ 133d68b`) with the KLOE-I cut file, card `nlo=nlo2=1`, `IFSNLO=1`, `ISR+INT+FSR`, `ivac=0`, `FF_pion=-1` ($F_\pi = 1$), no GVMD, no f_0 , soft cutoff $w = 10^{-5}$ (soft region $E_\gamma < w\sqrt{s}$). The production binary `phokhara_nostop` only disables the abort on “Max(i) too small” (Section 9); only weighted events are used.
- **BabaYaga** and the earlier PHOKHARA reference set (5 Aug) are taken from the Strong2020 comparison files.

Pipeline. SLURM arrays per run name; CEEX rank outputs are screened with `check_total.sh` and outliers are moved aside (`quarantine.sh`) before merging; `plot_compare.py` produces the distribution and ratio plots. Everything is in the repository (Appendix A).

3 The discrepancy

Table 1: Production runs (prod500k): PHOKHARA 300 seeds $\times$ 500k events; CEEX 29 tasks $\times$ 5 ranks $\times$ 500k at $E_{\min} = 10^{-5}$ GeV (one slow task cancelled).

	σ [nb]	$A(\theta^+)$
CEEX, $E_{\min} = 10^{-5}$ GeV	0.26836(4)	0.23752(14)
BabaYaga	0.268389(4)	0.23724(1)
PHOKHARA, this work	0.268353(46)	0.27872(16)
PHOKHARA, earlier reference	0.268393(23)	0.27861(8)

The new PHOKHARA run reproduces the earlier reference. All generators pass the charge-conjugation mirror test $\theta^+(\theta) = \theta^-(180^\circ - \theta)$, so this is not a mislabelled charge. Figure 1 shows the θ^+ distribution and Figure 2 the ratio of every observable to PHOKHARA: only the three pion angles are off, by a smooth tilt from 0.95 to 1.28.




Figure 1: Before the fix (prod500k): $d\sigma/d\theta^+$, CEEX/BabaYaga, and CEEX and BabaYaga over PHOKHARA (60 bins). CEEX ($E_{\min} = 10^{-5}$ GeV, black) and BabaYaga (blue) agree; PHOKHARA (red) is tilted.




Figure 2: Before the fix: every observable, CEEX and BabaYaga over PHOKHARA. Only θ^+ , θ^- and θ_{avg} are off; the spikes in $M_{\pi\pi\gamma}$ and M_{trk} are empty tail bins.

At LO the codes agree exactly (earlier runs with the Strong2020 form factor: $A_{\text{LO}} = 0.0779$ and $\sigma_{\text{LO}} = 2.6523$ nb in both), so the difference lies in the NLO correction to the charge-odd part of the cross section.

4 Strategy

The difference is charge-odd (C -odd): it flips sign under $\pi^+ \leftrightarrow \pi^-$ and integrates to zero over the symmetric cuts. At $\mathcal{O}(\alpha^4)$ the C -odd part of the cross section consists of the gauge-invariant classes $Q_e^5 Q_\pi^3$ and $Q_e^3 Q_\pi^5$ (the tree interference is $Q_e^3 Q_\pi^3$). Three tools were used to isolate it:

- **Cutoff independence.** The split into one-photon (Born + virtual + soft photons below the cutoff) and two-photon (hard) parts depends on the cutoff; the sum must not.
- **Matched split.** For the same cutoff each part is uniquely defined, so the two codes must agree part by part. PHOKHARA $w = 9.804 \times 10^{-6}$ ($w\sqrt{s} = 10^{-5}$ GeV) matches CEEX $E_{\min} = 10^{-5}$ GeV.
- **Momentum swap.** At a phase-space point, the C -odd part of any piece X is $X_{\text{odd}} = \frac{1}{2}[X(p_+, p_-) - X(p_-, p_+)]$. No charge parameter is needed.

Each histogram-level test used enough statistics to resolve $A(\theta^+)$ to $\sim 10^{-3}$ while checking that σ and $M_{\pi\pi}$ are unchanged. The plan with all tests and hypotheses is in `plan.txt`.

5 Tests

5.1 T1: soft-cutoff scan

Table 2: Unpatched PHOKHARA, 30 seeds $\times$ 100k per value of w .

w	10^{-3}	10^{-4}	10^{-5}	10^{-6}
$\sigma_{1\gamma}$ [nb]	0.1386	0.0955	0.0530	0.0103
$\sigma_{2\gamma}$ [nb]	0.1297	0.1728	0.2149	0.2580
σ [nb]	0.26833(33)	0.26827(35)	0.26792(33)	0.26830(30)
$A(\theta^+)$	0.3106(11)	0.2950(12)	0.2785(12)	0.2637(11)

$A(\theta^+)$ falls linearly in $\ln w$, by 0.0156 per decade, while σ is constant to 0.1%: the charge-even soft logarithms cancel between the one- and two-photon parts, a charge-odd one does not. In CEEX, $E_{\min} = 10^{-4}$ and 10^{-5} GeV give the same $A(\theta^+)$ to 10^{-4} .

5.2 T2: diagnostic variants

Three diagnostic PHOKHARA builds were run at $w = 10^{-4}, 10^{-5}, 10^{-6}$ (Figure 3a):

- **H1 fixed** (the parenthesis of Section 6.1): $A(\theta^+)$ moves by -0.037 and keeps its slope. H1 is a w -independent error.
- **B5 soft log removed** (probe): $A(\theta^+)$ becomes flat in w , but at 0.355. The whole w -dependence comes from this term; removing only the logarithm is not the right correction.
- **Both:** flat, at 0.319.




Figure 3: $A(\theta^+)$ against the PHOKHARA soft cutoff for the unpatched code, the T2 diagnostic variants and the fixed code (H1 + H3), with CEEX (grey band: $\pm 3\sigma$) and BabaYaga. The zoom (b) shows the fixed code at $30 \times 100k$, at $300 \times 500k$ (stars) and at $1000 \times 5M$ (diamond, drawn just left of $w = 10^{-5}$).


5.3 T3: split by photon multiplicity

CEEX-main was given an opt-in output of the per-multiplicity histograms (CEEX_NPH_HISTOS=1); PHOKHARA writes them with `print_summed=1`.

Table 3: Matched cutoff (10^{-5} GeV). CEEX 10 seeds $\times$ 5 ranks $\times$ 100k; PHOKHARA 30 $\times$ 100k (unpatched).

	σ [nb]		$A(\theta^+)$		angle shapes χ^2/ndf
	CEEX	PHOKHARA	CEEX	PHOKHARA	
1 photon (Born+V+S)	0.05268(7)	0.05239(24)	0.0375(13)	0.2360(44)	79–94
2 photons (hard)	0.21560(37)	0.21588(23)	0.2901(16)	0.2855(10)	1.0–1.4

The one-photon parts have the same σ and the same $M_{\pi\pi}$, θ_γ and E_γ shapes, but charge-odd cross sections of 0.0020 (CEEX) and 0.0124 nb (PHOKHARA). The difference, 0.0104 nb, is the whole total discrepancy (Figure 4). The hard two-photon parts agree.

**Figure 4:** T3: $d\sigma/d\theta^+$ of the one-photon (left) and two-photon (right) contributions at a matched cutoff, with the ratio PHOKHARA/CEEX.

5.4 T4: point-by-point comparison

988 accepted one-photon KLOE-LA points were dumped by a PHOKHARA variant (`phokhara_dump`) together with every piece of the one-photon matrix element, at the event and with the pion momenta swapped. A CEEX-main driver (`t4_ceex`) evaluates the tree, the hard virtual $\frac{\alpha}{4\pi}D$ (GoSam finite part plus mass counterterm), the wave-function counterterm and the YFS factor at the same momenta; the pion mass is set to PHOKHARA’s value. Each code is normalised to its own charge-even tree:

- tree C -odd/ C -even ratio: identical to 7×10^{-13} ;
- C -even NLO correction: CEEX -0.77566 , PHOKHARA -0.77570 , per-point difference rms 10^{-5} ;
- C -odd NLO correction: per-point difference D with rms 0.099;
- C -odd soft (eikonal/YFS) part: agrees, rms 7×10^{-4} .

D is exactly proportional to the tree interference at every point. Subtracting the two errors, each computed exactly per point ($v_s = 0.29655$, $v_f = 0.023\text{--}0.036$, $S_{\text{ISR}}(w) = -0.95556$),

$$\text{pred}_{\text{H1}} = B_4 - \frac{1}{2}B_3(v_f + v_s), \quad \text{pred}_{\text{H2}} = \frac{1}{2}B_3(1 + S_{\text{ISR}}(w)),$$

leaves a residual of rms 2.5×10^{-5} with mean 10^{-6} (Figure 5). H1 alone leaves 1.2×10^{-2} , H2 alone 8.8×10^{-2} .




Figure 5: T4, 988 points. Left: the C -odd difference D against the tree interference, with the H1+H2 prediction. Right: $\log_{10} |D - \text{subtracted}|$ for no subtraction, H1, H2 and H1+H2.

6 The two errors in PHOKHARA

The one-photon weight in the full-NLO mode is `amplit+helicityampL0+NLOpions+virt_pions`. In `ampL0` (inside `helicityampL0`) the pion branch adds, per photon polarisation: $B1 = |F|^2(1 + v_s)$, $B2 = |I|^2v_f$, $B3 = 2\text{Re}[I^*F]$ (tree interference), $B4$ and $B5$ below. Here I and F are the ISR and FSR tree amplitudes, v_s the electron-vertex correction at s and v_f the pion-vertex correction at q^2 , both $|M|^2$ -level factors with the analytic soft part removed.

6.1 H1: misplaced factor in the $\text{ISR} \times \text{FSR}$ correction

```

if((nlo2.eq.1).and.(pion.eq.1))then
  ampL0 = ampL0
  4    + dreal(
  5      dconjg(dme*(mb(1,1)-ma(1,1))*cvac_qq)*dme*(ddpl(1,1)-ddmi(1,1))*cvac_s    ! T1
  6      + dconjg(dme*(mb(2,2)-ma(2,2))*cvac_qq)*dme*(ddpl(2,2)-ddmi(2,2))*cvac_s    ! T2
  7      + dconjg((-ebppb*ma(2,1)+el_m2/ebppb*mb(2,1))*cvac_qq)
  8      *(-ebppb*ddmi(2,1)+el_m2/ebppb*ddpl(2,1))*cvac_s                        ! T3
  9      + dconjg((ebppb*mb(1,2)-el_m2/ebppb*ma(1,2))*cvac_qq)
  1     *(ebppb*ddpl(1,2)-el_m2/ebppb*ddmi(1,2))*cvac_s                        ! T4
  2     *(ver_f+ver_s) )                ! multiplies T4 only

```

(Continuation lines re-indented for reading; the code is unchanged.) The block is meant to be $\text{Re}[T_1 + T_2 + T_3 + T_4](v_f + v_s)$, the amplitude-level corrections on each side of $2\text{Re}[I^*F]$. As written it is $\text{Re}[T_1 + T_2 + T_3] + \text{Re}[T_4](v_f + v_s)$, so $\text{Re}[T_1 + T_2 + T_3]$ (about a third of the tree interference) is added with coefficient 1. The muon version of the same block (`phokhara_10.0.f:10384-10394`) is correct: `2*dreal(sum)*(ver_s+ver_f)/2`.

Fix (`variants/h1.patch`): `*cvac_s *(ver_f+ver_s) )` $\rightarrow$ `*cvac_s *(ver_f+ver_s).`

6.2 H2: tree and analytic ISR soft function in the ISR one-loop block

Block B5 contracts the FSR tree with the ISR current at one loop:

```
ampL0=ampL0+dreal( dconj(... ddpl, ddmi ...)*( ... apl, ami ... )    ! B5
```

The current (`hard_soft_part/ISRcurrent_no_mass_terms.f`) has the coefficient

```
V1(1)=2.d0*log(4.d0*w**2)*(log(m2)+1.d0) -4.d0*log(qq2)+3.d0*log(m2) + ...    ! 1.134
V1(1)=(1.d0-2.d0*alpha/4.d0/pi*V1(1))/Sp                                     ! 1.210
```

It therefore contains the tree (the “1”) and the analytic ISR soft function, and B5 carries an extra $\text{Re}[F^*I](1 + S_{\text{ISR}}(w))$ with

$$S_{\text{ISR}}(w) = \frac{\alpha}{\pi} \left[-\ln(4w^2) (1 + L) - \frac{1}{2}L^2 - L - \frac{\pi^2}{3} \right], \quad L = \ln \frac{m_e^2}{s}.$$

- The tree interference is already added by the `fsr=2` block (B3).
- In the full-NLO mode the analytic ISR soft function is replaced by the eikonal in `NLOptions`; it is subtracted from v_s , v_f and the leptonic tensor, but not here.
- At $w \simeq 10^{-5}$ the two parts nearly cancel ($1 + S_{\text{ISR}} = 0.044$), but S_{ISR} changes by 0.152 per decade of w . This gives the T1 slope: $\frac{1}{2} \cdot 0.152 \cdot A_{\text{LO}} \approx 0.019$ per decade, against 0.0156 measured.

Fix (variants/h3.patch): subtract $\text{Re}[\sum T](1 + S_{\text{ISR}})$ after B5, with S_{ISR} computed as in the v_s subtraction.

7 Result with the fix

`phokhara_fix` is the production binary with `h1.patch` and `h3.patch`.

Table 4: Fixed PHOKHARA against the soft cutoff.

w	10^{-3}	10^{-4}	10^{-5}	10^{-6}
<i>30 seeds $\times$ 100k</i>				
σ [nb]	0.26799(33)	0.26892(35)	0.26854(32)	0.26784(31)
$A(\theta^+)$	0.2349(12)	0.2352(13)	0.2383(12)	0.2374(11)
<i>300 seeds $\times$ 500k</i>				
σ [nb]	–	0.268345(48)	0.268351(46)	0.268354(44)
$A(\theta^+)$	–	0.23682(17)	0.23757(17)	0.23730(16)
<i>1000 seeds $\times$ 5M (5×10^9 events)</i>				
σ [nb]	–	–	0.268398(8)	–
$A(\theta^+)$	–	–	0.23725(3)	–

At high statistics the slope in w is -0.00023 ± 0.00012 per decade (unpatched: $+0.0156$), and the w -averaged asymmetries are

	$A(\theta^+)$	$-A(\theta^-)$	$-A(\theta_{\text{avg}})$
PHOKHARA, fixed (300 $\times$ 500k, w -average)	0.23724(10)	0.23722(10)	0.26651(9)
PHOKHARA, fixed (1000 $\times$ 5M, $w = 10^{-5}$)	0.23725(3)	0.23728(3)	0.26651(3)
CEEX, $E_{\text{min}} = 10^{-5}$ GeV (more1M, 148 $\times$ 5 $\times$ 1M)	0.23725(3)	0.23726(3)	0.26648(3)
BabaYaga	0.23724(1)	0.23727(1)	0.26650(1)

At $w = 10^{-5}$ and 10^{-6} the fixed PHOKHARA agrees with CEEX and BabaYaga within 0.3 – 1.5σ . The $w = 10^{-4}$ sample lies 0.0007 lower. To check it, the multiplicity split of Section 5.3 was repeated with the fixed code at a matched cutoff of 10^{-4} GeV (PHOKHARA $w = 9.804 \times 10^{-5}$, 100 seeds $\times$ 500k; CEEX $E_{\text{min}} = 10^{-4}$ GeV):

	CEEX	σ [nb] PHOKHARA (fixed)	CEEX	$A(\theta^+)$ PHOKHARA (fixed)
1 photon (Born+V+S)	0.09533(10)	0.09539(7)	0.1548(11)	0.1558(7)
2 photons (hard)	0.17290(23)	0.17301(5)	0.2826(13)	0.2829(2)
sum	0.26823(25)	0.26840(8)	0.2372(9)	0.2377(3)

Both parts agree, with angle-shape $\chi^2/\text{ndf} = 0.7\text{--}1.4$, and this sample gives $A(\theta^+) = 0.2377(3)$, consistent with $w = 10^{-5}$ and 10^{-6} . A 2% change of the cutoff cannot move $A(\theta^+)$ by 0.0009, so the low $w = 10^{-4}$ value is a fluctuation of that sample (2.6σ from this one). T4 at $w = 10^{-4}$ also leaves no systematic C -odd residual after H1+H2 (mean -7×10^{-6}).

Highest statistics. A final run of 1000 seeds $\times$ 5M events (5×10^9 events passing the cuts, $w = 10^{-5}$) gives $\sigma = 0.268398(8)$ nb and $A(\theta^+) = 0.23725(3)$, $-A(\theta^-) = 0.23728(3)$, $-A(\theta_{\text{avg}}) = 0.26651(3)$. To match this precision on the CEEX side, a further CEEX run (more1M: 5 MPI ranks $\times$ 1M events per task at $E_{\text{min}} = 10^{-5}$ GeV, seeds 101–300, independent of prod500k) was started; the 148 tasks finished at the time of writing (732 rank outputs; 8 whose total diverged were discarded by `check_total.sh`) give $\sigma = 0.268389(9)$ nb and

$$A(\theta^+) = 0.23725(3), \quad -A(\theta^-) = 0.23726(3), \quad -A(\theta_{\text{avg}}) = 0.26648(3),$$

in agreement with the fixed PHOKHARA and with BabaYaga to 3×10^{-5} in all three asymmetries. Figures 6 and 7 repeat Figures 1 and 2 with these two runs: the pion-angle shape χ^2/ndf of CEEX against the fixed PHOKHARA drops from ~ 1900 to 0.8–1.4, and against BabaYaga it is 0.9–1.2.




Figure 6: After the fix: as Figure 1, with the fixed PHOKHARA (1000 $\times$ 5M, $w = 10^{-5}$) as the denominator and CEEX more1M (148 $\times$ 5 $\times$ 1M, $E_{\text{min}} = 10^{-5}$ GeV). The tilt is gone.




Figure 7: After the fix: every observable, CEEX (more1M, $E_{\min} = 10^{-5}$ GeV) and BabaYaga over the fixed PHOKHARA (1000×5M).

8 Caveats and open points

- Only the configuration studied here was tested ($F_\pi = 1$, no VP, KLOE-LA cuts). Both errors are on code paths active for `pion=1` with `nlo2=1` in any configuration, so they should also affect the Strong2020 form-factor results (there the NLO gap in $A(\theta^+)$ is 0.012: CEEX 0.0804, PHOKHARA 0.0924). This has not been re-run.
- After the fix, one of the three high-statistics samples ($w = 10^{-4}$) lies 0.0007 below the others; an independent sample at the same cutoff agrees with them and with CEEX part by part (Section 7), so it is treated as a fluctuation.
- Muons take a different, correct branch and are not affected.
- Totals and charge-even distributions are not affected by either error (σ unchanged in every run; charge-even per-point agreement 10^{-5}).
- The 0.00027 (2σ) excess of CEEX prod500k over BabaYaga in $A(\theta^+)$ was a statistical fluctuation: CEEX more1M gives $A(\theta^+) = 0.23725(3)$, equal to BabaYaga (0.23724(1)) and the fixed PHOKHARA (0.23725(3)). more1M was still running when this report was written (148 of 200 tasks); the numbers will be updated with the full run.
- With 5×10^9 PHOKHARA events, small BabaYaga–PHOKHARA shape differences become visible in a few charge-even observables (χ^2/ndf 2–3 in $p_\pm$, $p_\pm^\perp$, E_γ ; 10 in $|p_-^z|$, where

BabaYaga’s binning already gave ~ 20 before). CEEX agrees with PHOKHARA there at $\chi^2/\text{ndf} \approx 1$. They are unrelated to the pion-angle issue.

- Tests T5 (charge-class split) and T6 (two-photon point-by-point) of the plan were not needed: the two-photon part agrees (T3) and H1+H2 explain the one-photon part (T4).

9 Other findings

- The checked-out PHOKARA_10.0 has an empty stub `variables_cuts.f` (no cuts, no histograms); the KLOE-I cut file comes from the Strong2020 scenario repository. The card format and the FF_pion numbering changed with respect to the Strong2020 cards.
- PHOKHARA aborts without writing histograms when a weight exceeds the scanned maximum (`phokhara_10.0.f:1910,2135,2263`). Dropping such seeds, as the Strong2020 merge did, biases the result; for weighted events the abort is unnecessary.
- `virt_pions` overwrites its quad-precision momentum arguments (~ 1 GeV). This is harmless in production but breaks any second call in the same event.
- In the 5 Aug reference folder, `Phokhara/NL0novp_*.csv` is the muon result (0.9925 nb), not the pion one.
- CEEX: a few rank outputs per hundred diverge at 100k events (11/300); they are removed before merging. The CEEX-main README still documents `-ref_alloc` and `-vfreeze`, which main no longer parses.
- Cost: at equal precision on σ (prod100k), PHOKHARA needed about 45 times less CPU than CEEX: 0.25 against 28 CPU-hours.

A Reproduction and files

<code>plan.txt</code>	test plan (T0–T6, hypotheses H1–H5)
<code>RUN/config.sh</code>	all paths
<code>RUN/RUN_CEEEX/</code> , <code>RUN/RUN_PHOKHARA/</code>	SLURM submitters and job scripts, runcards
<code>RUN/RUN_PHOKHARA/variants/</code>	<code>h1.patch</code> , <code>h3.patch</code> (the fix), <code>h2.patch</code> (T2 probe), <code>dump.patch</code> (T4), <code>build.sh</code>
<code>RUN/RUN_CEEEX/t4/</code>	CEEX point driver; <code>MC_nphsplit.patch</code> (per-multiplicity histograms)
<code>RUN/POST_PROCESS/</code>	<code>merge</code> , <code>outlier check</code> , <code>phok_table.py</code> , <code>t3_compare.py</code> , <code>t4_compare.py</code>
<code>RESULTS/</code>	<code>prod500k</code> , <code>t1_t2</code> , <code>t3_split</code> , <code>t4_points</code> , <code>t4_fix</code> , <code>fix500k</code> , <code>final500k</code> , <code>fix5M</code> , <code>final5M</code> , <code>final5M_c1M_e5</code> , <code>prod500k_e5</code>
<code>REPORT/</code>	this report, <code>make_figures.py</code> , <code>make_numbers.py</code>

Building the fixed binary:

```
cd RUN/RUN_PHOKHARA/variants
./build.sh fix h1.patch h3.patch          # -> src/phokhara_fix
cd ../..
EXE=RUN_PHOKHARA/src/phokhara_fix \
  RUN_PHOKHARA/submit.sh <run> <n_seeds> <nges> 10000 <walltime>
```